



IMT Atlantique

Bretagne-Pays de la Loire
École Mines-Télécom

ELASCRIFT: A DSL FOR CODING ELASTICITY IN CLOUD COMPUTING

Simon Dupont (*)

Salma Bouri, Frederico Alvares, Thomas Ledoux (**)

(*) *Sigma Informatique*

(**) *IMT Atlantique campus Nantes, LS2N, INRIA*

OUTLINE

1. PROBLEM

- 1.1 Cloud Elasticity
- 1.2 ElaStuff Overview

2. ELASCRIP

- 2.1 Overview
- 2.2 Examples
- 2.3 Guarantees

3. VALIDATION

- 3.1 Testbed
- 3.2 Results and discussion



IMT Atlantique
Bretagne-Pays de la Loire
École Mines-Télécom

Overview

Adjust resources automatically according to the demand so as to satisfy a certain level of Quality of Service (QoS) while minimizing infrastructure costs

Mainly based on IaaS elasticity (VM scaling)

Shortcomings

Some limits in terms of responsiveness (e.g., VM startup time)

Risks in terms of

- ✓ Over-provisioning: highly cost
- ✓ Under-provisioning: SLA violation


Why?

A problem of reactivity


A problem of resource granularity



Objective: software layer (SaaS) may take part in the global elasticity process to cope with IaaS elasticity shortcomings



Elasticity @ IaaS Level

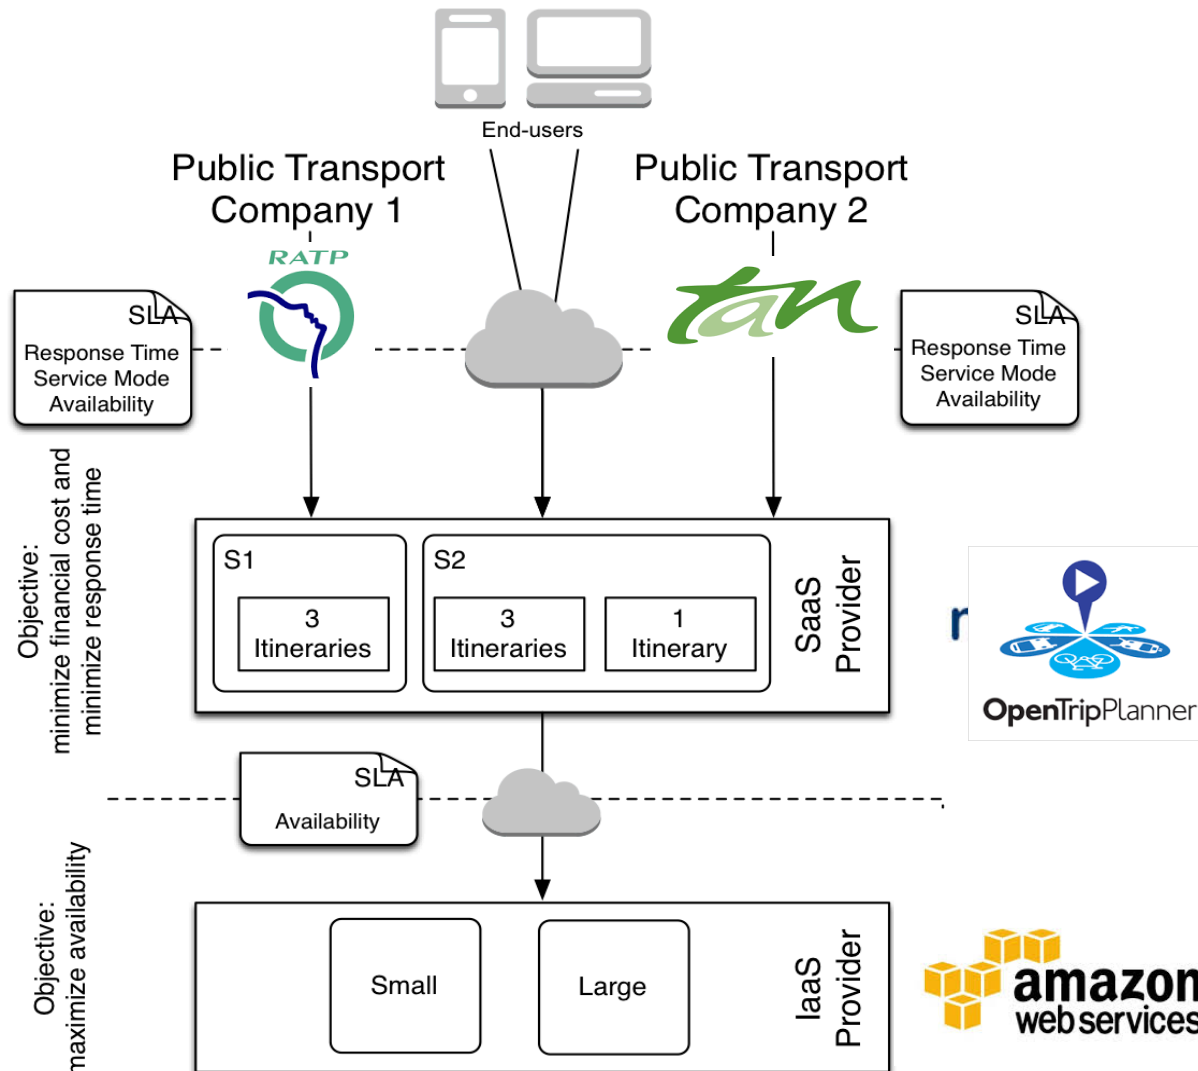


Elasticity @ SaaS Level

CROSS-LAYER ELASTICITY

Illustration

5



Conclusion

Advantages

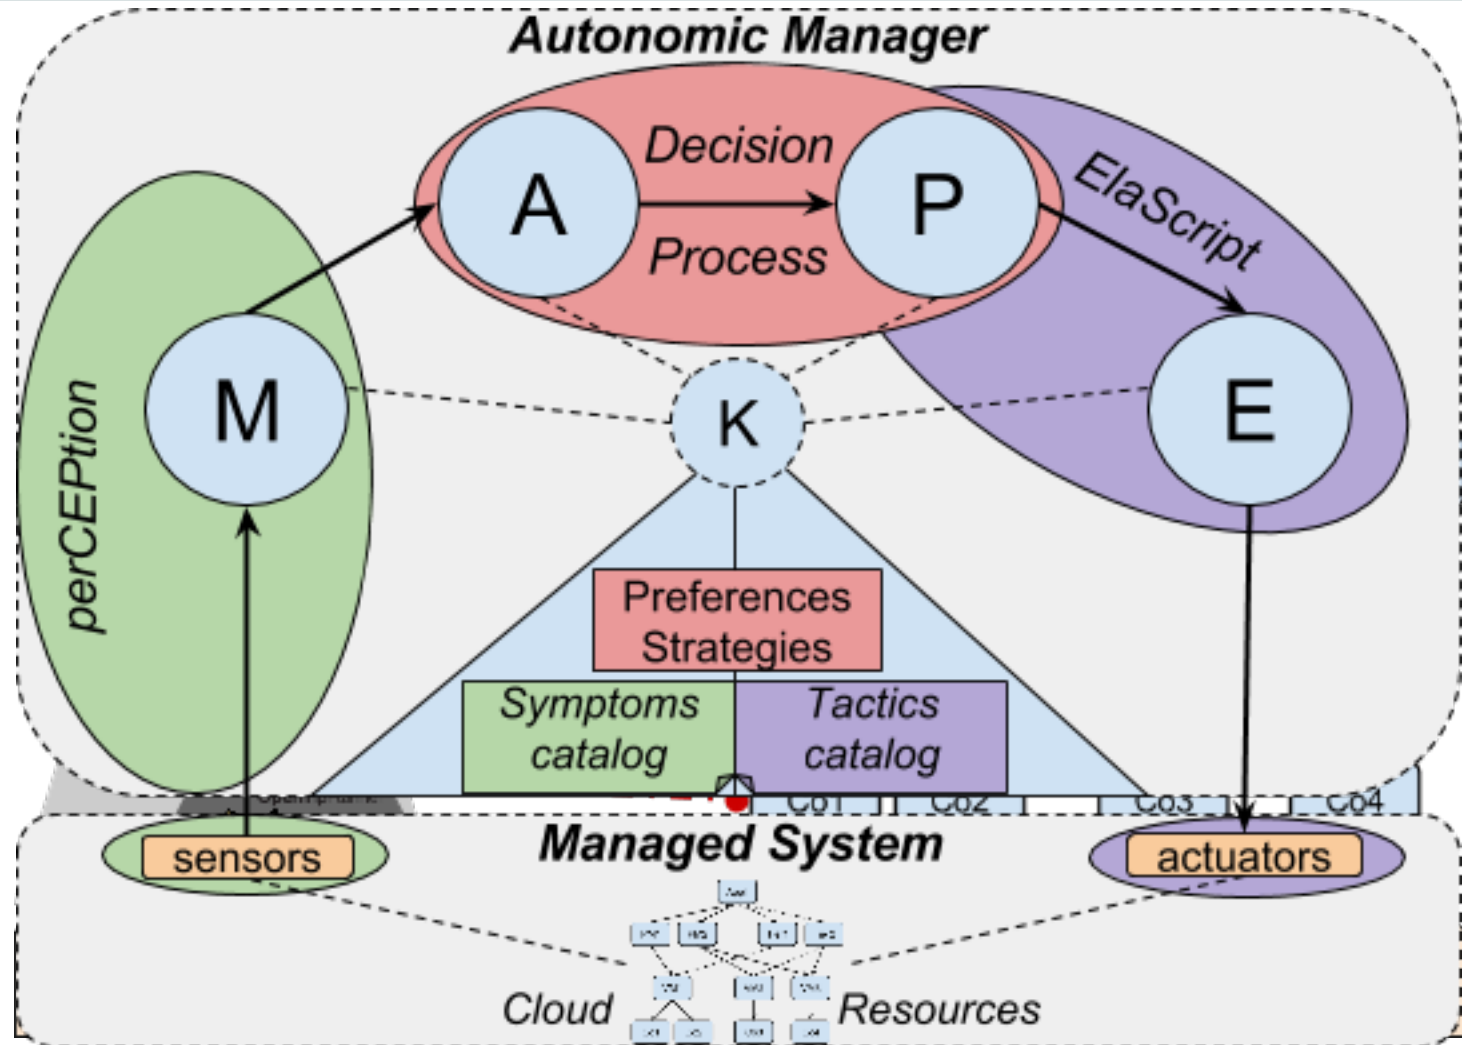
A finer scaling granularity

Fulfill SLA contracts (more reactivity)

Minimize resources (economic and ecological pros)

Consequence

Cloud administrators can create more sophisticated reconfiguration plans and strategies (e.g., minimize energy consumption)



Goal

A support for Cloud administrators to simply and safely express reconfiguration plans


How ?

Manipulate resource graphs via 8 primitive APIs

- ✓ (Horizontal, vertical scaling) X (IaaS, SaaS)

Use a DSL (Domain Specific Language) to wire and update the graph

- ✓ Possibilities for static analysis, verification or optimization
- ✓ DSL constructs for coordination of cross-layer elasticity



A scripting language

Imperative, Java-like syntax, static typing, no new variables, only one kind of iterator

Guarantees

Naming, scoping, typing, business rules

Technologies used

EMF, Xtext, Xtend, Java framework Executor

tactic {
 begin
 when SymptomName (Type1 resource1 , . . . , TypeN resourceN)
 do
 resource1.actionX;
 . . .
 resourceN.actionY;
 end

Simple example

```
begin
  when TierHighRT (Tier myTier)
  do myTier.soi ( "medium");
end
```

Cross-layer example

```
begin
  when Tier VMsOverloaded (Tier tier, List<VM> vms)
  do
    [
      tier.soi (2 , "small ") ;
      ||
      foreach vm in vms do { vm.sds;}
    ]
    foreach vm in vms do {vm.sus; }
  end
```

Scoping rules

```
1 begin
2 when TierOverloaded(tier)
3 do
4     // For each vm of the tier
5     foreach vm in tier do {
6         // ScaleUpInfra
7         vm.sui;
8     }
9     vm.sui;
10
11 end
12
13
```

⚠ Couldn't resolve reference to Variable 'vm'.
1 quick fix available:
➡ [Change to 'tier'](#)
Press 'F2' for focus

Typing rules

```
1 begin
2 when ComponentHighRT(comp)
3 do
4     comp.sos("myNewComp", 4);
5
6 end
```

⚠ SOS action can't be executed on Component resource

Business rules

```
1 begin
2 when VmLowCpu(vm)
3 do
4 [
5     vm.sii;
6     vm.sdi(1);
7     ||
8     vm.sis("myComp");
9 ]
10 end
```

vm is locked: can't execute another action on this resource after SII action.

Press 'F2' for focus

Business rules

```
1 begin
2 when TierHighRT_VmsOverloaded(tier, vms)
3 do
4     // Execute in parallel
5     [
6         // Add 2 small instances to the tier
7         tier.soi(2, "small");
8         // Degrade the offering soft of one level for all overloaded vms's components
9         foreach vm in vms do {
10             vm.sds;
11         }
12     ]
13     // Switch back (upgrade) the offering soft to nominal state
14     foreach vm in vms do {
15         vm.sus;
16     }
17 end
```

The SOI action takes time, perhaps you can take the opportunity to parallelize other actions during that time...

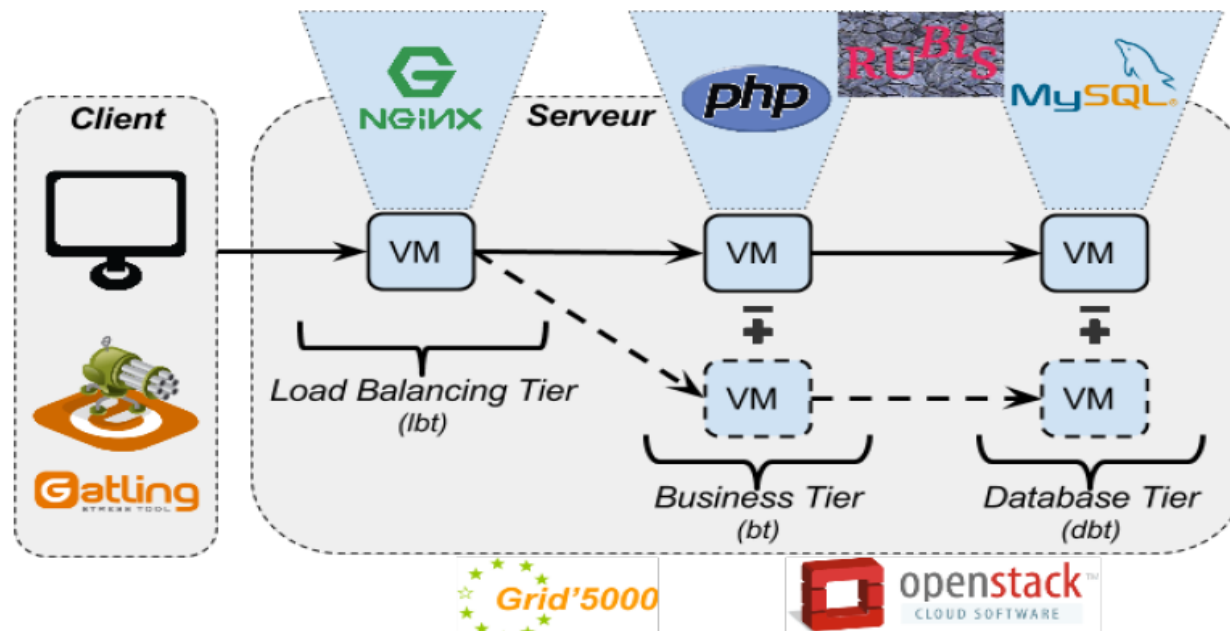
SaaS application

RUBiS (an eBay like auction site)

A three-tier architecture

Three Software offerings for recommendation:

- ✓ zero, one (user-to-user) or two (user-to-user and item-to-item)



IaaS infrastructure

OpenStack on Grid'5000 (large-scale testbed for experiment-driven research)

2 physical machines with wattmeter

Workload Traces

Real trace of FIFA'98

Elastuff Tactic

```
begin
  when App[Low | Normal | High]RT (App application)
  do application.sds( [2 | 1 | 0]);
end
```

Experiment Runs

Base line: a static version of RUBiS with 2 recommendations vs Elastuff one

QoS Criteria

Availability

- ✓ Baseline requests failed x330 more

Performance

- ✓ Baseline response time x8

QoE

- ✓ Elastic version: 18% with 0 recommendation, 10% with only one

Energy consumption

- ✓ equivalent

Run	Request Succeeded	Requests Failed	Requests served in each Off_{soft}			95th percentile Response Time (ms)	Power Consum. (W)
			noReco	oneReco	twoReco		
Baseline	1124449	25922	0	0	255114	7890	3416
Elastic	1149926	78	52352	29765	205261	895	3280

SLA contracts

Trade-off between criteria are indirectly managed by ElaScript

Conclusion

ElaScript: a DSL that enables Cloud administrators to simply and concisely program complex multi-layered reconfiguration plans

Perspectives

Consider new kinds of Cloud resources

- ✓ software containers (e.g., Docker), microservices

Consider new dimensions/actions

- ✓ VM migration

Renewable energy and Software elasticity

- ✓ Software offering should be relied on the presence of green resources
→ improve the carbon footprint and the software at the same time

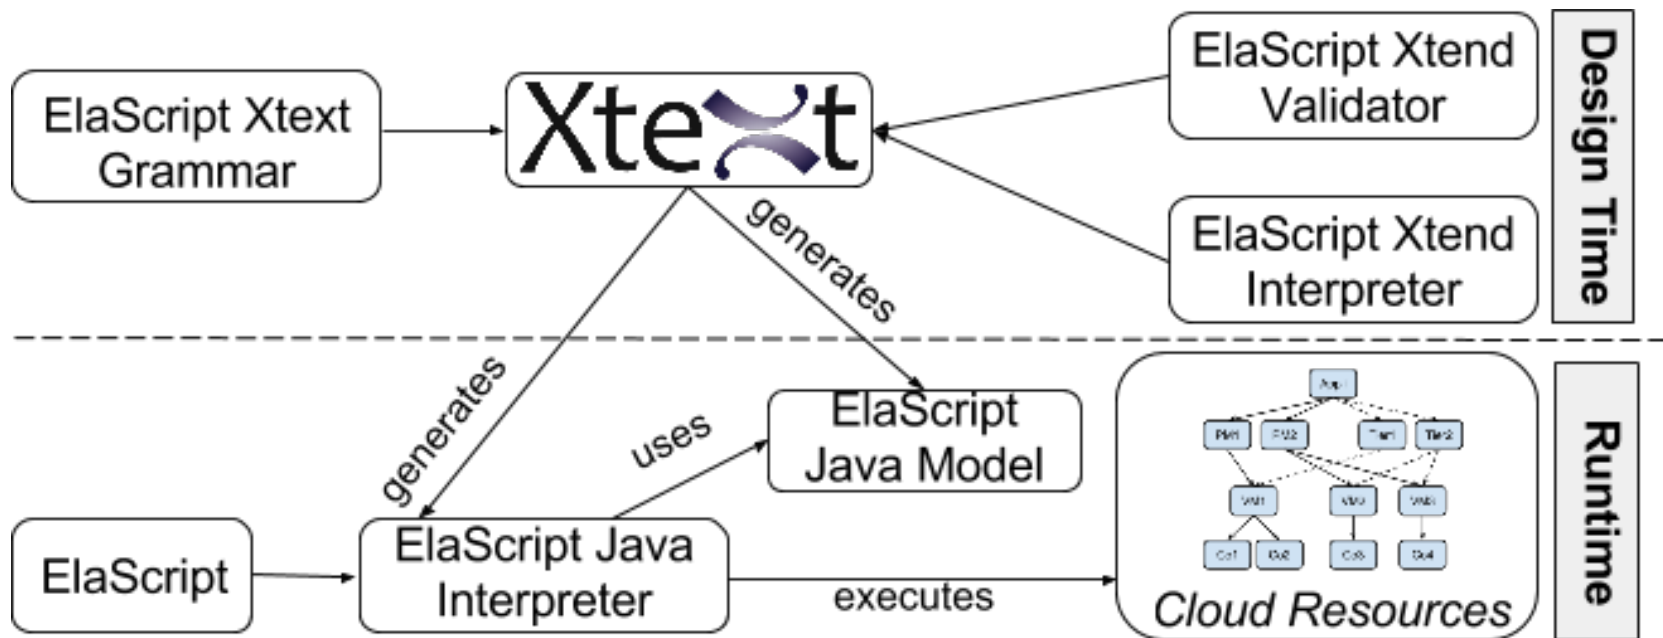
Towards a *Elasticity-as-a-Service* offering

THANKS!

Any questions?



IMT Atlantique
Bretagne-Pays de la Loire
École Mines-Télécom



ELASTUFF ECOSYSTEM

How to define the resource graph?

```
OfferingCO mode1
OfferingCO mode2
OfferingCO mode3
```

```
OfferingVM small { vcpu 1 ram 1024 disk 10 cost 47 }
OfferingVM medium { vcpu 2 ram 2048 disk 50 cost 94 }
OfferingVM large { vcpu 4 ram 4096 disk 500 cost 188 }
```

```
typePM medium { cpu 4 ram 4096 disk 500 }
```

```
VM VM1 , small { CO CO1 , MySQL , mode1 ; CO CO2 , MySQL , mode2 }
VM VM2 , small { CO CO3 , MySQL , mode1 }
VM VM3 , small { CO CO4 , MySQL , mode1 }
```

```
Application App1 { [ Tier Tier1 { VM1}; Tier Tier2 { VM2 ; VM3} ]
[ PM PM1 , medium { VM1 } ; PM PM2 , medium { VM2 ; VM3} ]
}
```

